

Lecture 06 – Processes and Threads

Jaejin Lee

Dept. of Computer Science and Engineering, College of Engineering

Dept. of Data Science, Graduate School of Data Science

Seoul National University

<http://aces.snu.ac.kr>

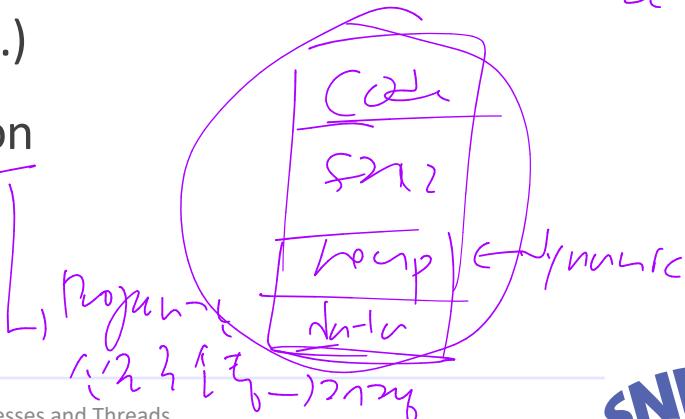
- A process is an instance of a computer program that is being executed.

- A stream of instructions being executed
- Abstraction used by the operating system

- A process consists of:

- Registers (ALU, Registers) → 32bit 주소 지정 R1 ~ R16 or 32
- Memory (code, data, stack, heap, etc.) → 8비트 단위로 주소 지정
- I/O status (open file tables, etc.)
- Signal management information

OS 가 두 개의 프로그램을 실행하는 경우



Ex. 웹툰, 인터넷 브라우저...
Same Program or
CPU?



Supervisor Mode vs. User Mode

- Modern processors provides two different modes of execution:

- Supervisor (kernel) mode**

- All instructions can be executed in supervisor mode (also known as protected mode, system mode, monitor mode, or privileged mode)

- For the operating system kernel

- User mode**

- The processor is allowed to execute only a subset of the instructions
- For all other software (including the remaining part of the operating system) than the kernel

- Example:**

- I/O instructions** are privileged instructions

- An application needs to request an I/O service to the operating system to perform I/O operations

System Calls

- A system call is the way how a program running in user mode requests a service to the operating system
 - Typically implemented with a trap (a.k.a. an exception or a fault)
 - A trap instruction invoked by the program triggers a trap, resulting in a switch to kernel mode
 - The kernel performs some action to handle the trap before returning control to the program

trap이 걸리면
CPU가 커널로
 okre인해준다



Uniprogramming vs. Multiprogramming

- Uniprogramming

1 프로그램

2 프로그램 실행.

(1 프로그램 + 흑백 출력).

- Only one process at a time
 - DOS
- Poor resource utilization

- Multiprogramming

- Multiple processes at a time
 - Modern operating systems, such as Windows, Unix, Linux, etc.
- Increases resource utilization

수행 중인 프로세스:

1. P1 → 10ms
2. P2 → 10ms

or
Core
CPU

멀티 프로그램 + 퍼포먼스

비교적 짧은 시간에 다양한 작업 수행.

Virtual Memory

- The operating system's abstraction of the physical memory in the system

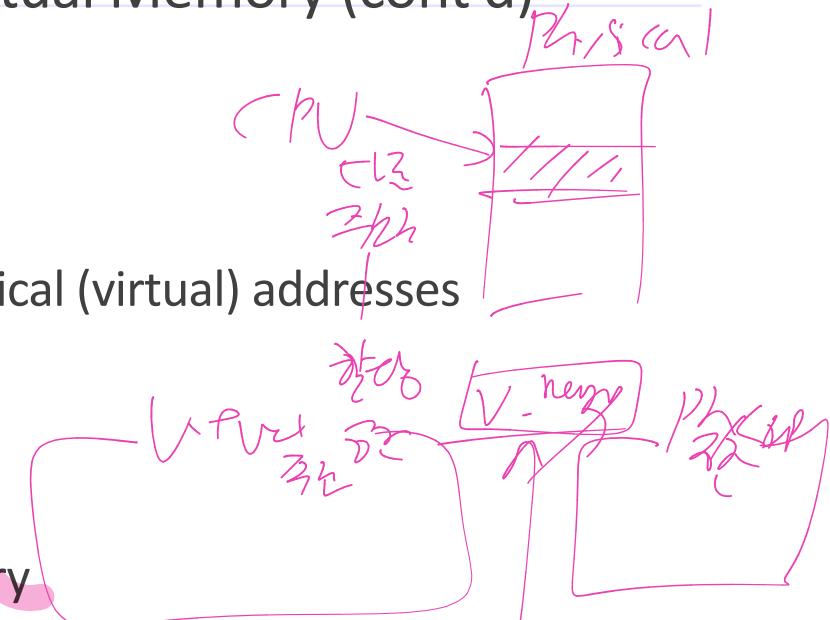
(→ 운영체제가 물리 메모리 A₁, A₂ 를 가상 메모리 A₁, A₂로 만든다.)

- Provides each process with the illusion that the process has ~~a local~~ exclusive use of the memory and a much larger memory space than that available in the system



Virtual Memory (cont'd)

- Logical (virtual) address
 - An address generated by the CPU
 - Logical address space - the set of all logical (virtual) addresses generated by a program
- Physical address
 - An address seen by the physical memory
 - Physical address space - the set of all physical addresses corresponding to the logical addresses
- The virtual memory in the operating system is in charge of the run-time mapping from virtual to physical addresses
 - Exploits a hardware device called the memory-management unit (MMU) for fast translation between virtual and physical addresses



MMU



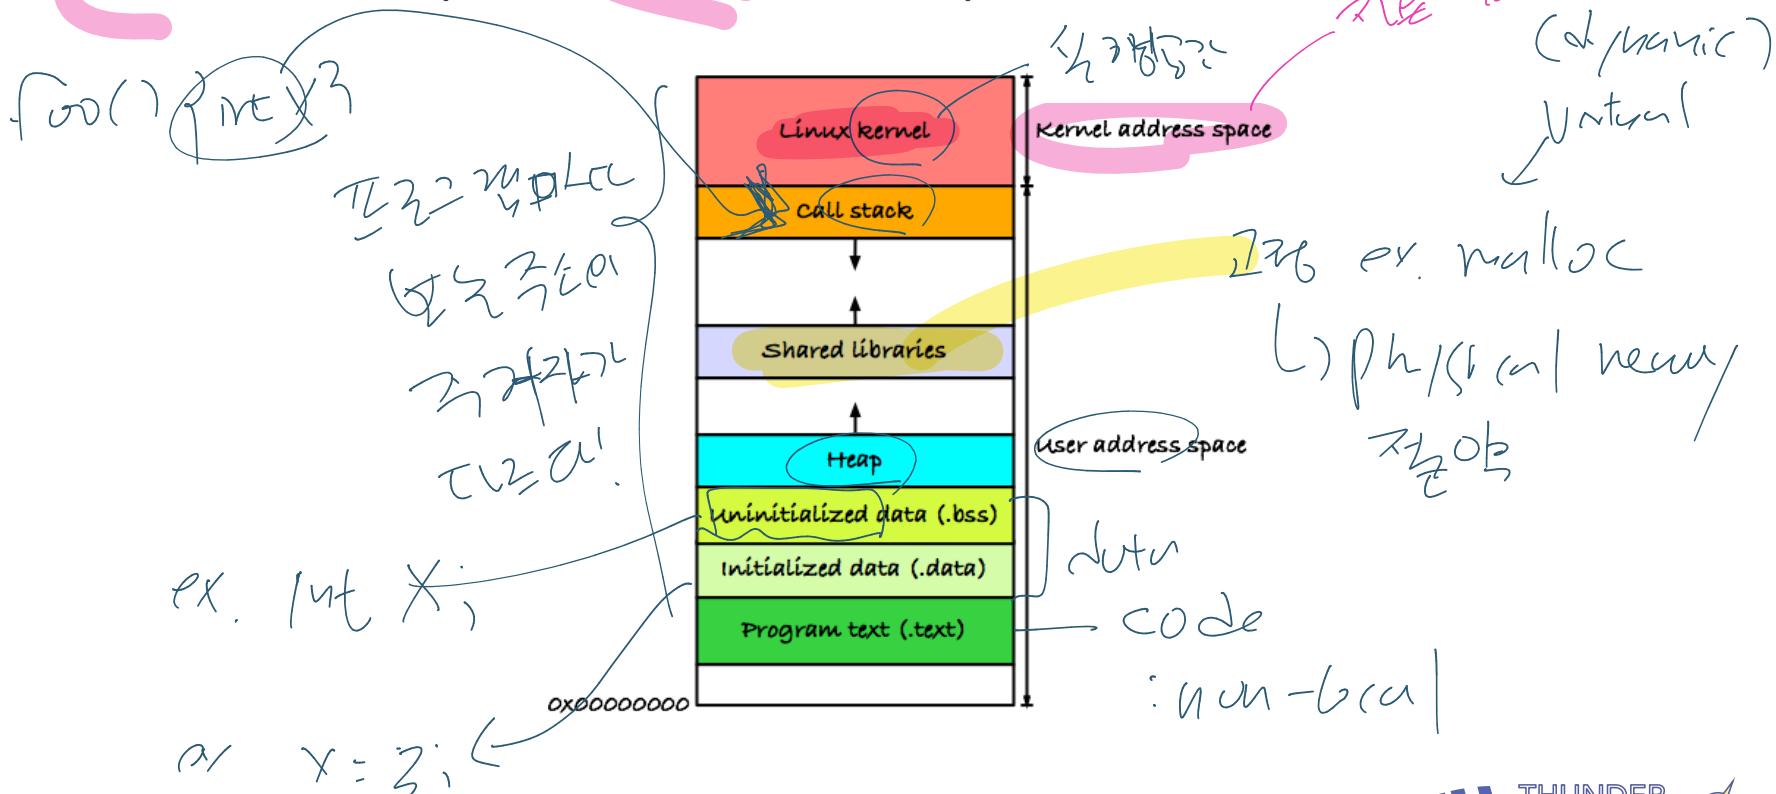
THUNDER
Research Group
서울대학교 천동 연구실



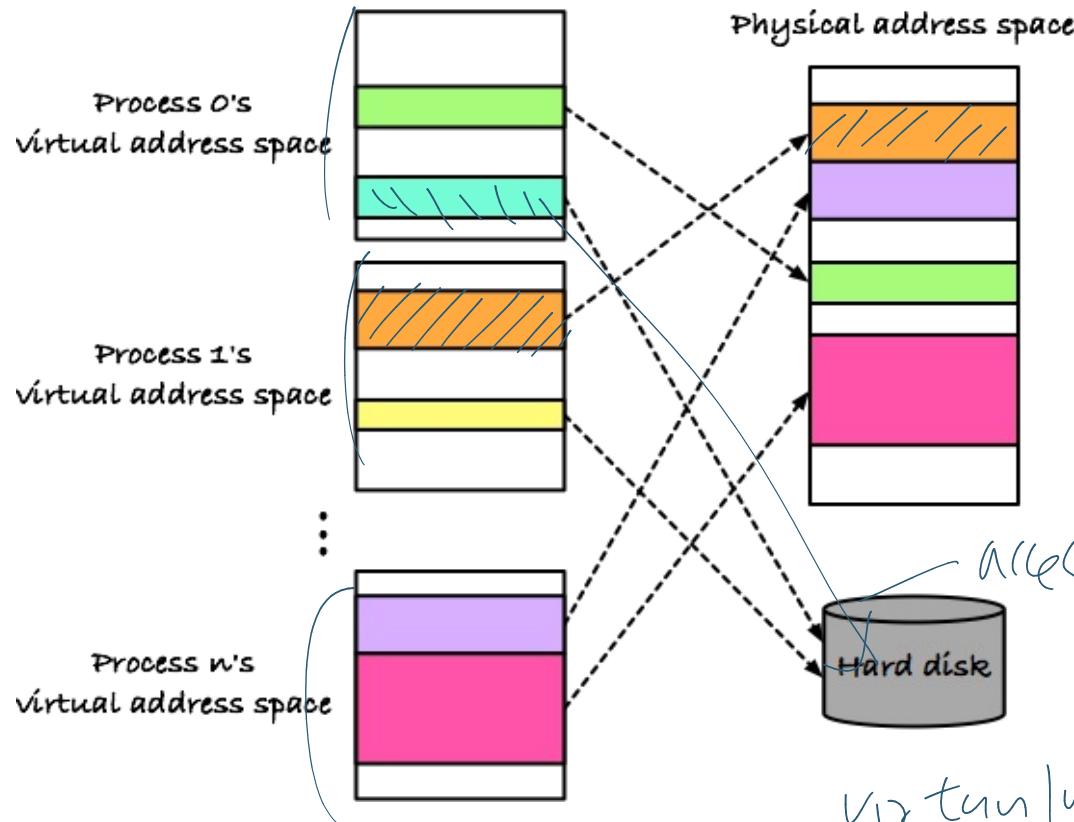
Process - 3/27-1

Address Space of a Process

- A process has its own private address space (virtual address space)
 - A process cannot affect the state of another process directly
 - Memory protection
- Kernel address space vs. user address space



Address Space of a Process (cont'd)



Processes → Process (进程)

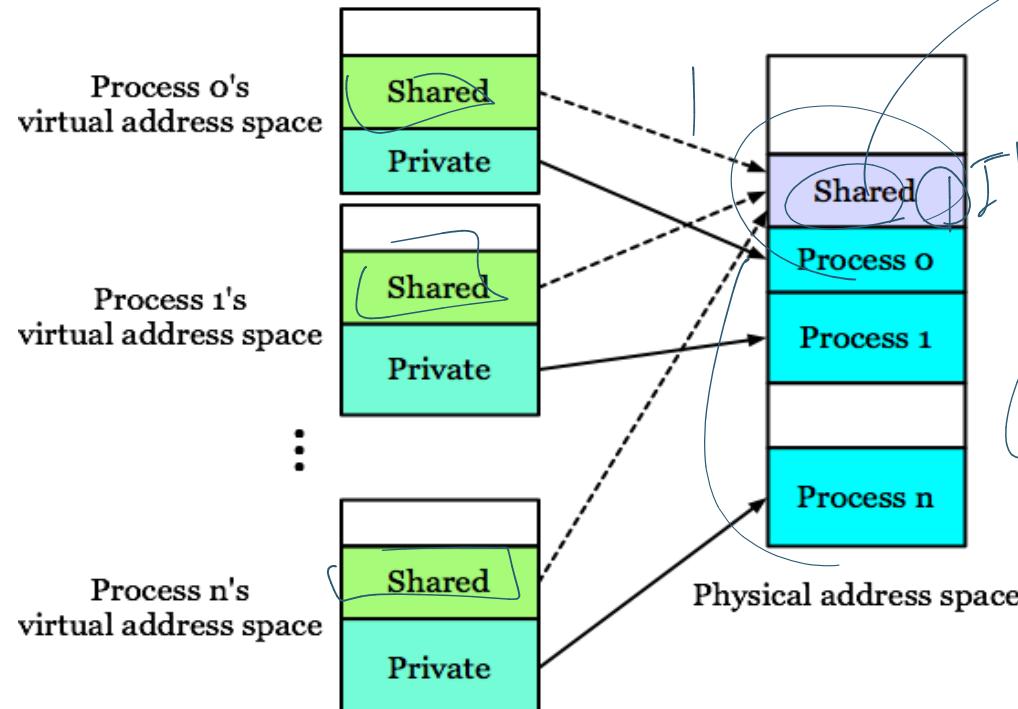
Communication Between Processes

- Cooperation and coordination between processes are accomplished by writing and reading to a location in the shared address space
- Another way to achieve them is using an interprocess communication (IPC) mechanism
 - The IPC is a way of exchanging data between processes without sharing any portion of their virtual address space
 - Expensive

=>

IPC ①

Communication Between Processes (cont'd)



进程与线程 (CPU多路复用)

Concurrency

- A computer system is typically a multiprogramming system

- Kernel processes execute system code

- User processes execute user code

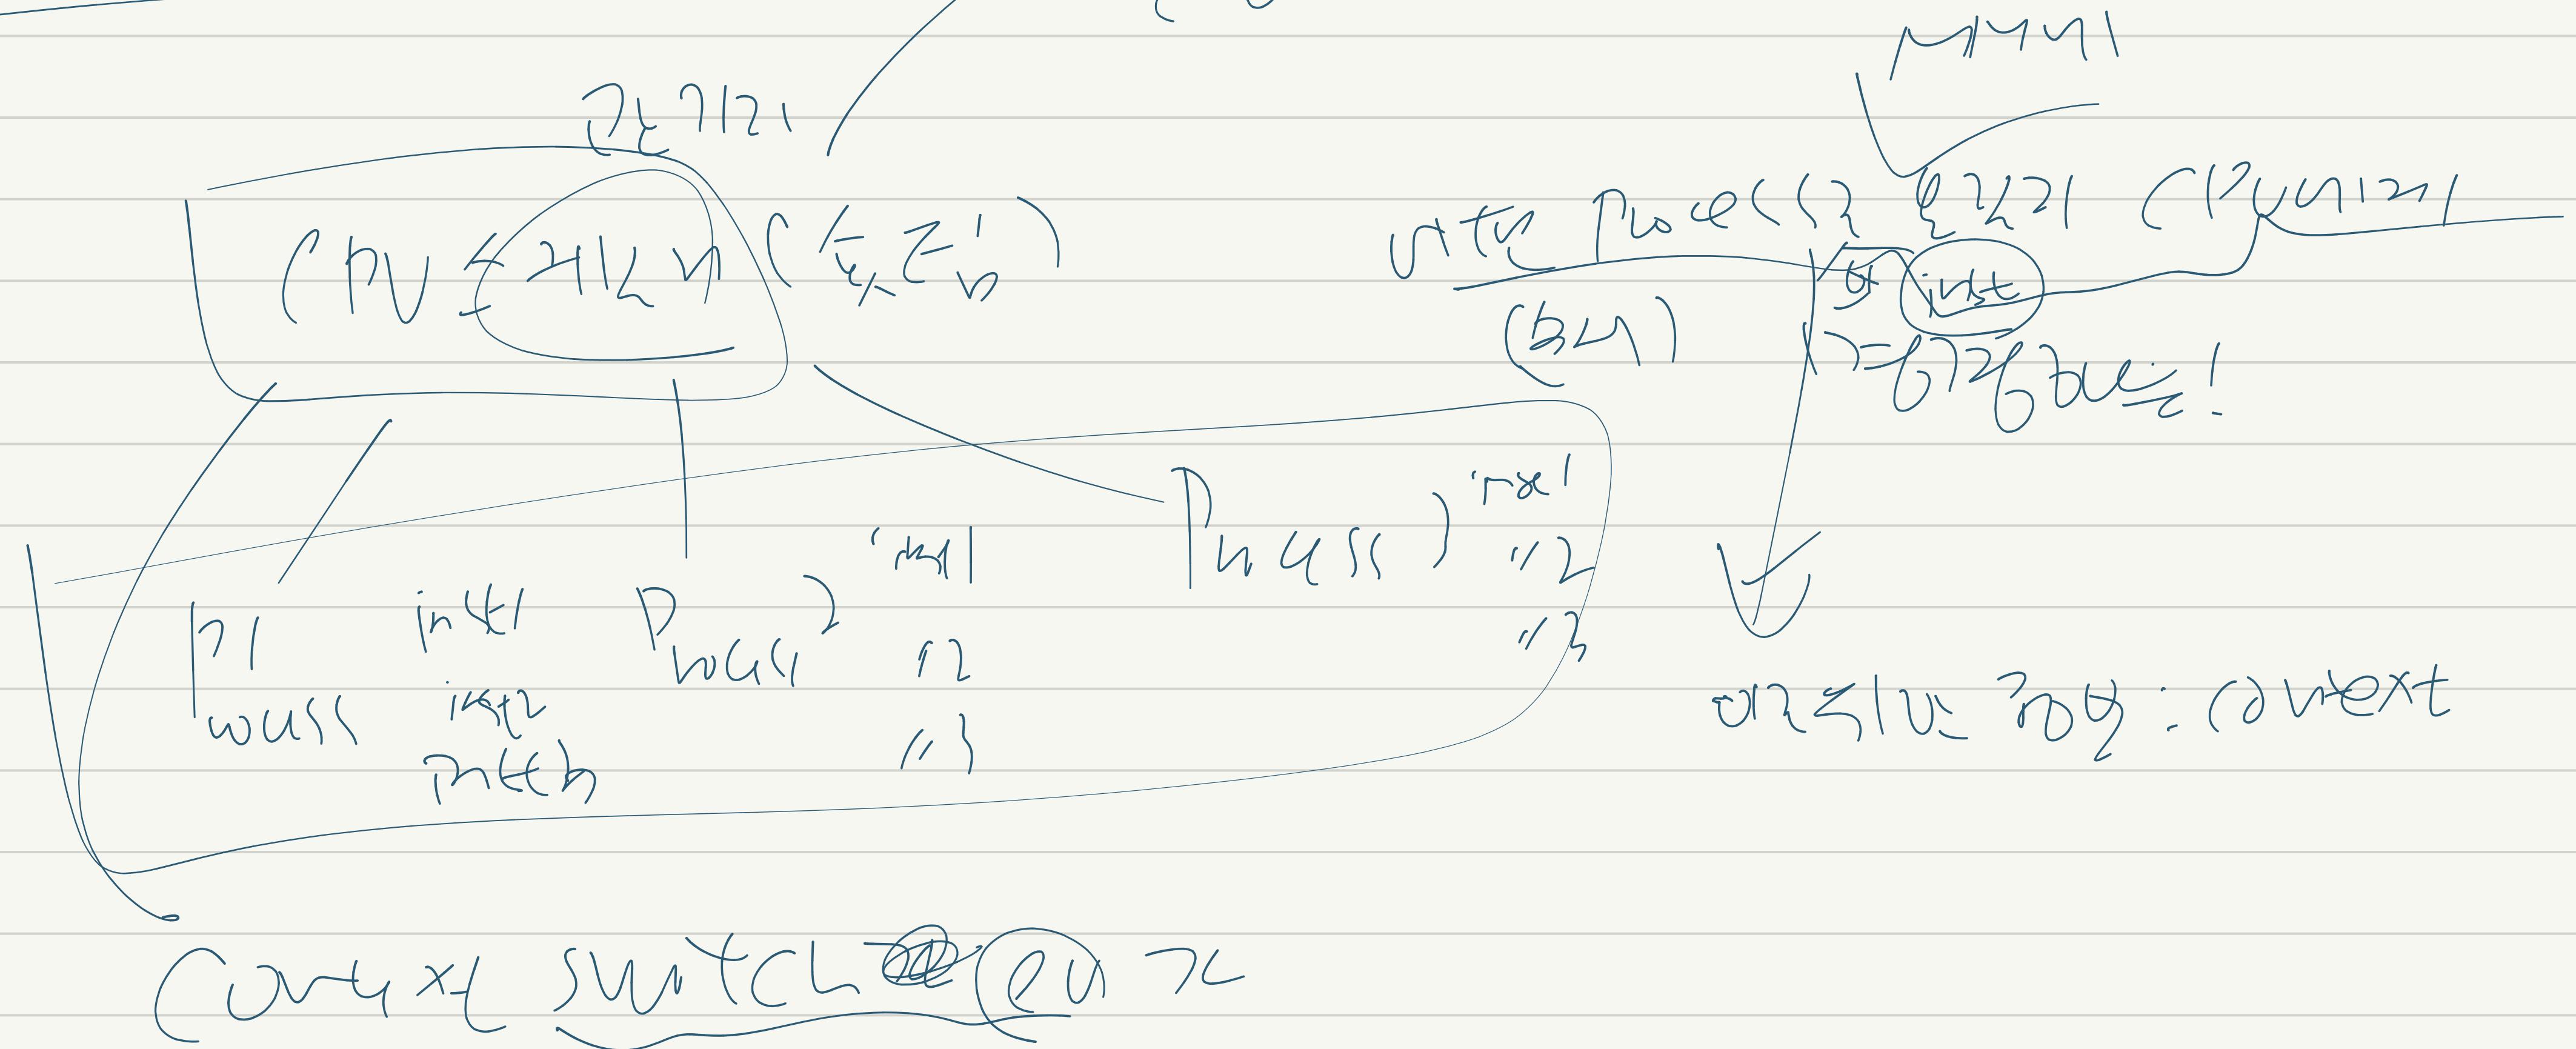
- All these processes may execute concurrently on the same system

- Concurrency
 - Instructions of one process are interleaved with those of another process
 - Implemented by context switches

cf. Concurrency vs Parallel Process
=> 共享 + 独立资源
=> 并行

什么时候并行比串行好?

→ CPU Exchange mit der Bus, was blieb? Context switcher



* Context switch im M³? Scheduler?

Context Switch

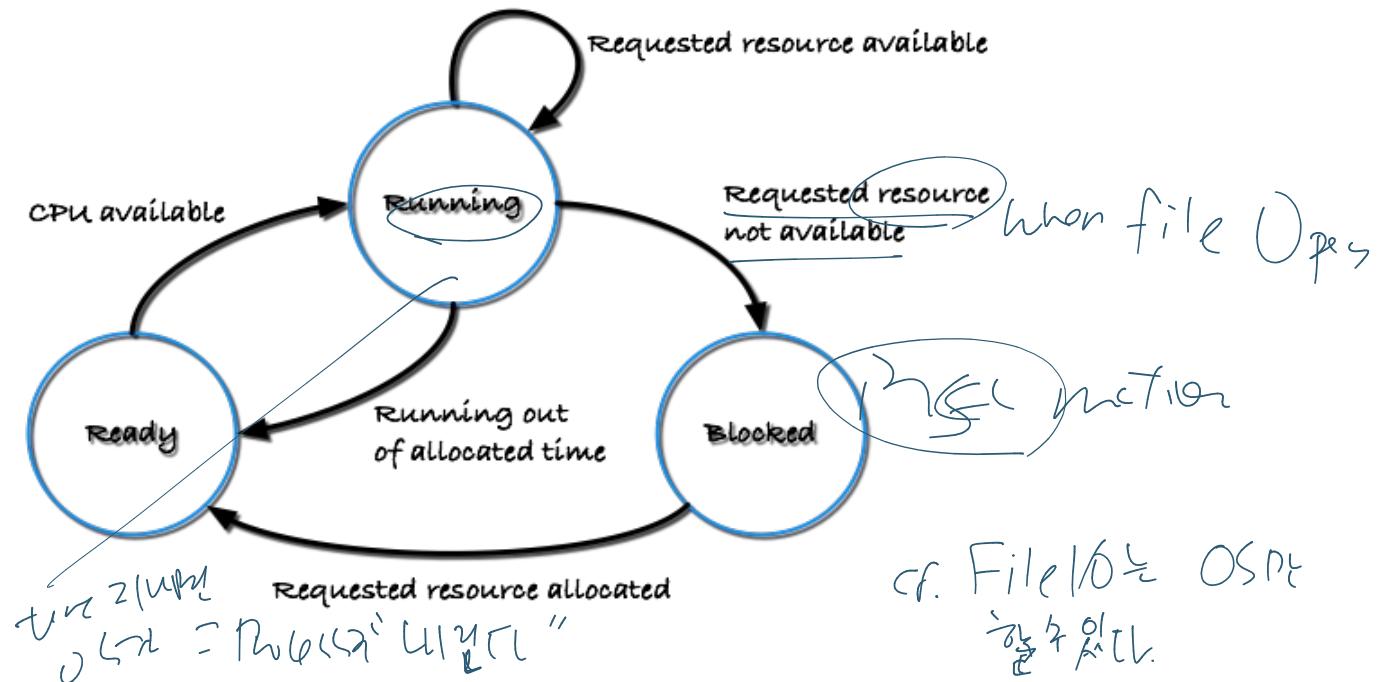
- The CPU switches back and forth from process to process to execute the instructions from different processes
 - The CPU scheduler in the operating system is in charge of it
- A process context is the information that must be saved before a context switch occurs to allow the continuation of the process later
- The operating system transfers control from the current process (say, p) to another process (say, q) after saving the context of p and restoring the context of q
- Then, control is passed to the location of q 's code where it left off due to a previous context switch



Process State Transitions

- When a program runs, the corresponding process changes state
 - Running: using the CPU
 - Ready: no CPU available
 - Blocked: waiting for some event (e.g., I/O) to occur

↳ OS가 실행을 허용하는 경우



2021년 7월 2일

Preemptive vs. Cooperative

Preemptive multitasking

- Permits preemption of tasks
- All processes will get some amount of CPU time at any given time
- More reliably guarantee each process a regular slice of operating time
- Nearly all modern operating systems support preemptive multitasking

설명

Cooperative multitasking

- Tasks must be explicitly programmed to yield when they do not need system resources (e.g., CPU)
- Rarely used in these days

OS가 Task(P)를 주는

내가 주고자, (rand when or

Priority)

ex. nice

yield

X GIL

join or wakeup?

Threads

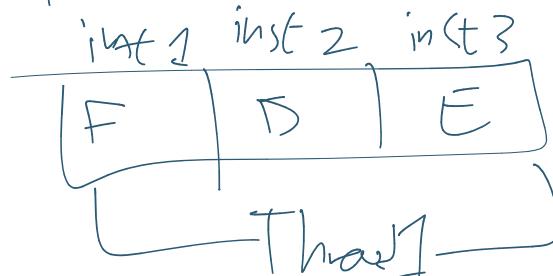
Thread of control

- Independent Fetch/Decode/Execute loop
- The smallest unit of processing that can be scheduled by an operating system

CPU 5 21329 ei 템스

- A thread logically consists of:

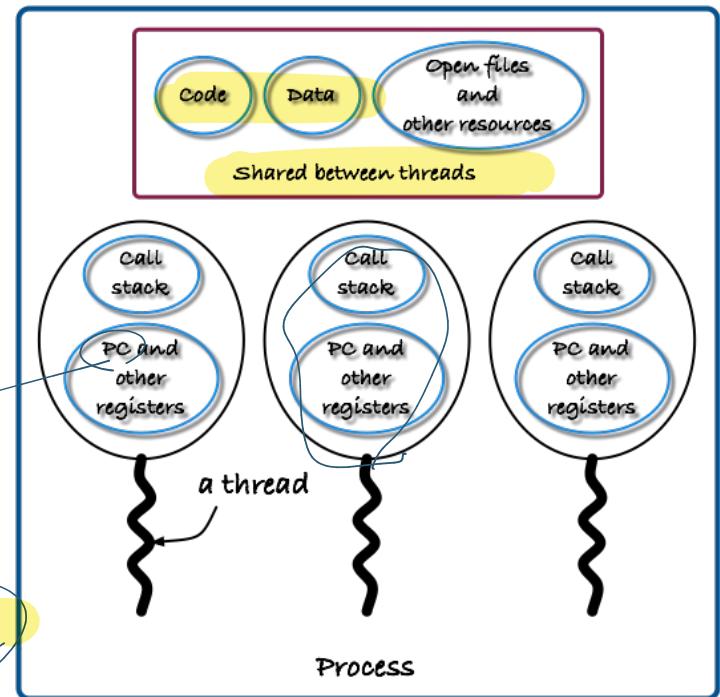
- Code
- Registers
- Stack
- Thread-local data



User-level thread vs. kernel-level thread

Threads (cont'd)

- In general, a thread is contained in a process
 - Multiple thread can exist within the same process
 - Share resources with other threads
 - Code
 - Data
 - OS resources: open files, signals, etc.



Communication Between Threads

- Multiple threads within a process share portions of the virtual address space of the process (e.g., text (code) and data sections) by default

Threads in the same process

Share memory!

(进程) 内存共享
进程间通信

- Cooperation and coordination between threads in the same process is accomplished by reading and writing variables allocated in the shared space

- Writes to a shared address by one thread is visible to reads of the other threads

→ IPC

Thread Library

- Provides the programmer an API for creating and managing threads
 - A user-level library entirely in user space with no kernel support
 - A kernel-level library supported directly by the operating system
 - Code and data structures for the library exist in kernel space
 - An API function call typically results in a system call

POSIX Pthreads

- A user-level library

进程 | Process

p : POSIX

Pthreads

Linux Schedulers

- Completely Fair Scheduler (CFS)
 - Since kernel 2.6.23
- No distinction between processes and threads in scheduling
- To maintain fairness in providing processor time to processes
- A run-queue for each processor
 - Contains processes whose state is 'ready'
- Nice values
 - A processes' relative weight used in CFS
 - Lower nice value → higher weight → higher priority

threads & process in Linux

multi-tasking on Linux (OS)

idle process

process in memory

process in mutual

The quantum mechanical
of Energy levels and Time Slices

- Time slice

- The time interval for which a process can run without being preempted
 - Proportional to the processes' weight

processes weight
Priority $\uparrow \rightarrow$ time \downarrow \rightarrow \uparrow
High priority

- Virtual runtime

- A measure for the amount of time provided to a given process
 - The smaller a processes' virtual runtime, the higher its need for the processor

Processor

↳ 하지 않거나 하는 시스템 있는가?
Realtime은 높은 priority + ?
↳ CPU 캐시 =) 처리速度快
Processor가 더 빠르면

2 processor
~~12012~~
CPU -> 1202
CPU 대신에
CPU 대신에

Physical task \leftarrow

Virtual Runtime

- A processes' cumulative execution time inversely scaled by its weight
 - The weight is a decay factor for the time for which a process has run

Priority $\uparrow \rightarrow$ Virtual runtime \downarrow .

$$\text{virtual runtime}(P_i, t) = \frac{W_0}{W_{P_i}} \times \text{physical runtime}(P_i, t)$$

W_0 = the weight for the nice value 0

(weight \propto priority) \propto Virtual Runtime

Red-black Tree

97"

- A self-balanced tree
- No path in the tree will ever be more than twice as long as any other
b h(x) bound
- Operations on the tree occur in $O(\log n)$, where n is the number of nodes in the tree

Red-black Tree (cont'd)

- CFS maintains a red-black tree ordered by the virtual runtime
 - Run-queue
 - Maintained independently for each processor
 - The process with lowest virtual runtime is the left-most leaf node (highest: the right-most leaf node)
 - The scheduler picks the left-most node to schedule next to maintain fairness
 - The task will be added to the tree with a new virtual runtime after running
- ex.
Processors
MPI Processors
virtual runtime
- the most
runny (35)

- Performs scheduling on each **scheduling tick**
- Decrement the **time slice** of the currently running process P by the **tick period**
 - When the time slice reaches 0, a flag is set
- Update the virtual runtime of P
- Check the flag
 - If set, preempt P and insert it to the **run-queue**
 - Schedule the process in the **left-most node** in the red-black tree

1-2022년
Update
7/13
CFS Algorithm

252
by 20ms per tick.
CPU调度 = RR⁴
Tick ex. 2.3ms
Preempt
as process
physical
time.

multiple processor

SMP Scheduling

symmetric Multiple Processing

- CFS

- Scheduling processes for a single processor

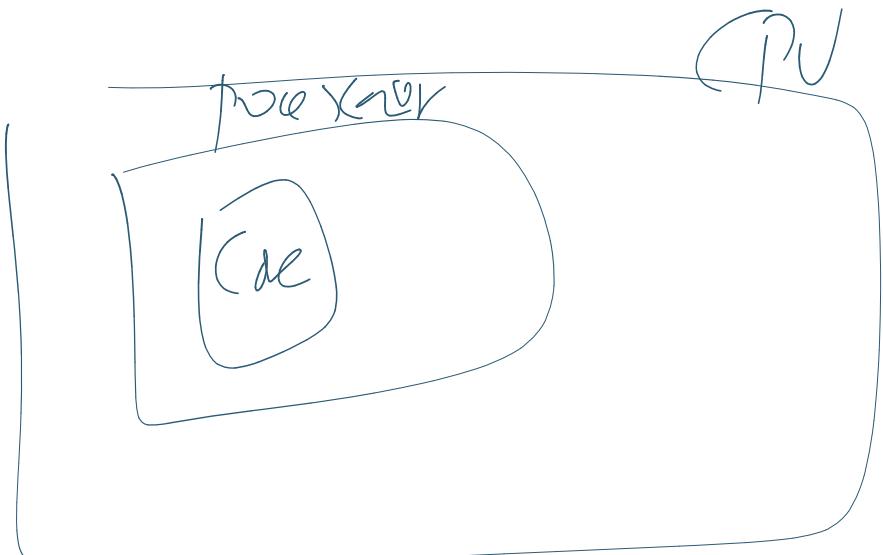
- Run-queue load balancing

- Distribute processes across multiple processors

(Ranking)

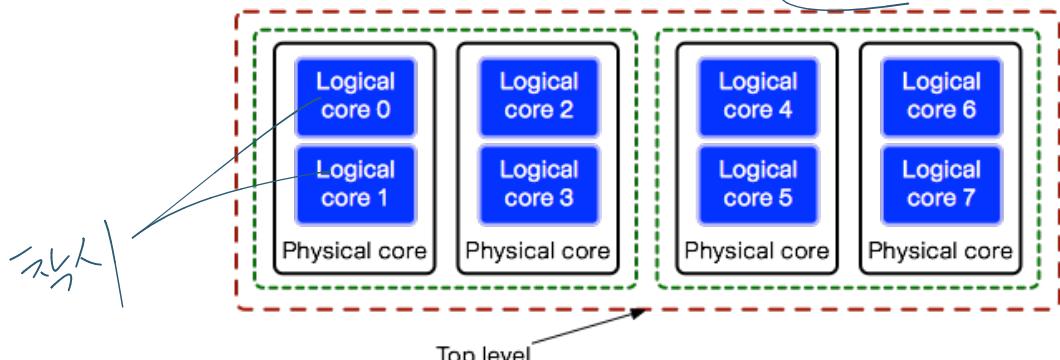
0113021 process 5
↑ 21121011044

의 경우에
balance
by user
각각



Scheduling Domains

- A set of processors whose workloads should be kept balanced by the kernel
 - Share properties and scheduling policies
 - Can be balanced against each other
 - Partitioned in one or more groups
 - Hierarchically organized
 - Top scheduling domain: the set of all processors in the system
 - Each scheduling domain contains policy information which controls how decisions are made at that level of the hierarchy
 - An active load balancing is executed regularly which moves up the scheduling domain hierarchy and checks all groups along the way if they got out of balance
 - If so it does a balancing attempt considering the policy rules of the corresponding domain



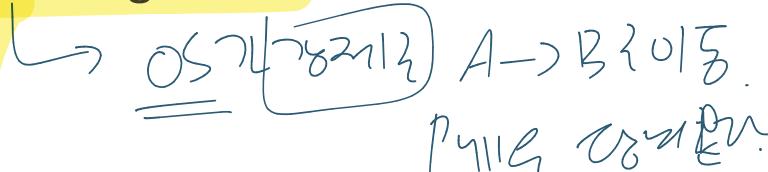
Run-queue Balancing

Excerpt from T3

- Perform load balancing on each rebalancing tick

주로

- Push migration



- Check hierarchically if a scheduling domain is significantly unbalanced

processor 마다 짐加重

- Find the busiest run-queue in the domain

- By calculating the load of each processor or group

Load of a processor: run-queue length

Processor 짐加重 계산!

- Migrate processes from the busiest run-queue to another one

cf. halting problem

(유저) (설정)
Push Pull
push pull

Push vs. Pull

Cache Use (설정)

↓ due to because
contain

- Push migration

- A specific process periodically checks the load on each processor and evenly distributes the load by moving (or pushing) processes from overloaded to idle or less-busy processors

- Pull migration

- Occurs when an idle processor pulls a waiting process from a busy processor

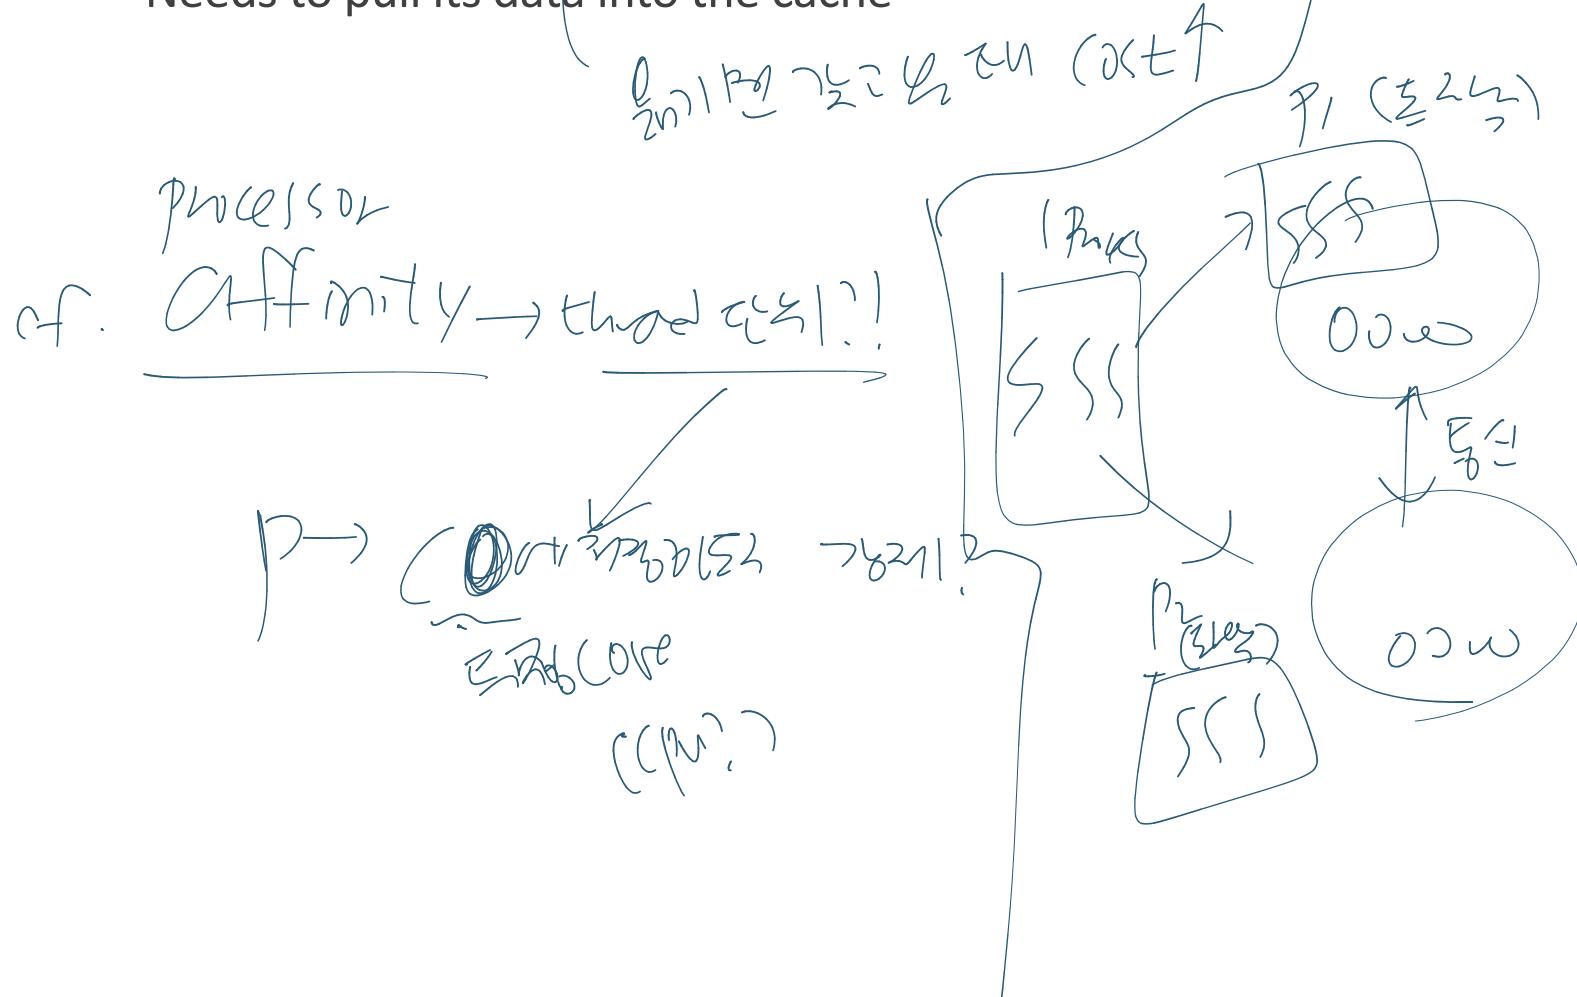
Cache : idle migration

- The Linux scheduler implements both techniques

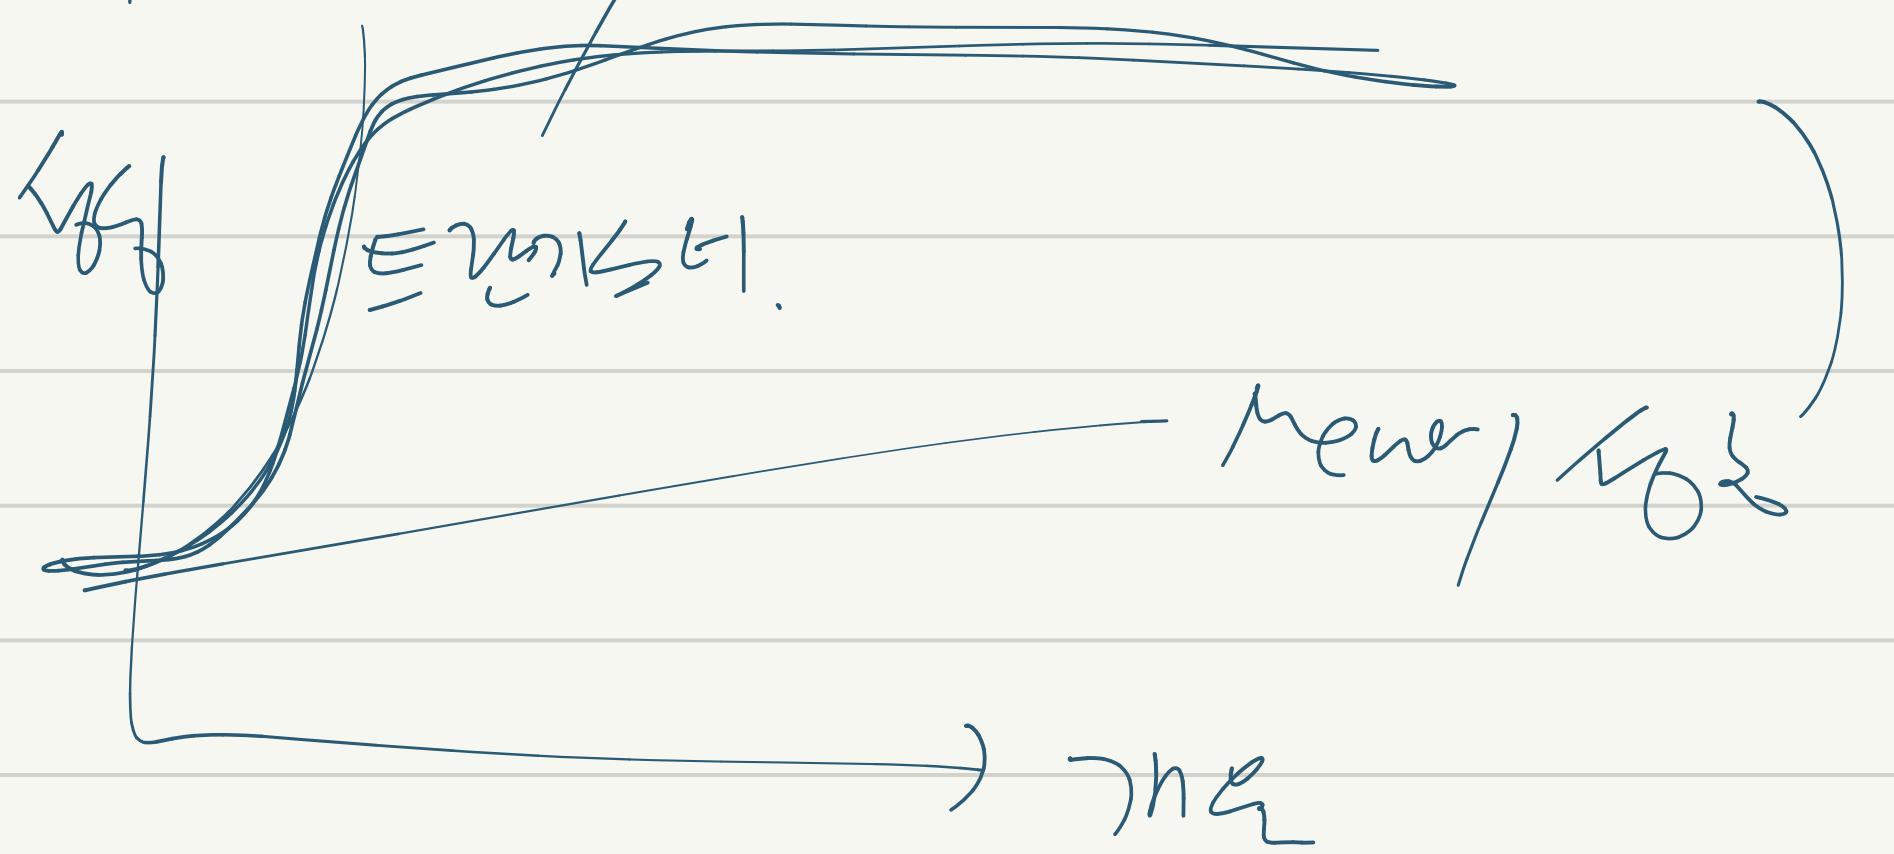
- Linux runs its load balancing algorithm every 200 milliseconds (push migration) or whenever the run-queue for a processor is empty (pull migration)

Negative Aspect of Process Migration

- The new processor's cache is cold for a migrated task
 - Needs to pull its data into the cache



Cf. Memory wall



Cache miss
Solve

Cf. Single core vs. 쌍핵 이상,
=> 메모리 벽

